A Regularization Approach for Prediction of Edges and Node Features in Dynamic Graphs ¹

Abstract

We consider the two problems of predicting links in a dynamic graph sequence and predicting functions defined at each node of the graph. In many applications, the solution of one problem is useful for solving the other. Indeed, if these functions reflect node features, then they are related through the graph structure. In this paper, we formulate a hybrid approach that simultaneously learns the structure of the graph and predicts the values of the node-related functions. Our approach is based on the optimization of a joint regularization objective. We empirically test the benefits of the proposed method with both synthetic and real data. The results indicate that joint regularization improves prediction performance over the graph evolution and the node features.

1 Introduction

Forecasting the behavior of systems with multiple responses has been a challenging problem in the context of many applications such as collaborative filtering, financial markets, or bioinformatics, where responses may be, respectively, movie ratings, stock prices, or activity of genes within a cell. Statistical modeling techniques have been widely applied for learning multivariate time series either in the multiple linear regression setting [3] or with autoregressive models [19]. More recently, kernel-based regularized methods have been developed for multitask learning [7, 2]. These approaches share in common the use of the correlation structure between input variables to enhance prediction of every single output. Frequently, the correlation structure is assumed to be given or is estimated separately. A discrete encoding

Emile Richard : ENS Cachan - UniverSud - CMLA, France, 1000mercis, France, emile.richard@cmla.ens-cachan.fr

Andreas Argyriou : TTI Chicago, argyriou@ttic.edu

Theodoros Evgeniou : Technology Management and Decision Sciences, INSEAD Bd de Constance, Fontainebleau 77300, France, theodoros.evgeniou@insead.edu

Nicolas Vayatis : ENS Cachan - UniverSud - CMLA, UMR CNRS 8536 France

nicolas.vayatis@cmla.ens-cachan.fr

of correlations between variables can be modeled as a graph so that learning the dependence structure amounts to performing graph inference through the discovery of unobserved edges on the graph. The latter problem is interesting per se and is known as link prediction, where it is assumed that only part of the graph is actually observed [10, 8]. This situation occurs in various applications such as recommender systems, social networks, or proteomics, and the appropriate tools can be found among matrix completion techniques [18, 5, 1]. In the realistic setup of a time-evolving graph, matrix completion was also used and adapted to take into account the dynamics of the features of the graph [14].

In this paper, our goal is to simultaneously predict multiple outputs defined over the vertices of a time-evolving graph and learn the structure of the graph. One important assumption we make is that the network effect is a cause and a symptom at the same time. Consider, for instance, the problem of forecasting sales based on purchase data from an e-commerce market where the semantic information is not reliable. The data can be represented as the (bipartite) graph of users and products connected by purchases which evolves over time according to the purchase history. We expect that within a cluster of co-purchased items, the sales should be correlated, and reversely, correlated sales should induce more edges between corresponding items. A similar situation arises in the context of financial data where the graph reflects dependencies between stocks and one aims at both predicting stock prices and inferring the dependence structure. We build on the approach proposed in [6] which shows that, in the case where graph structure information is available about the dependencies between input variables, the efficiency of multivariate predictions can be enhanced. In this paper, we tackle a problem of broader scope, that is, we assume that the underlying graph is unknown. We propose a formulation of the problem as a regularized risk minimization which balances both objectives of prediction and matrix completion at the same time. The proposed formulation is convex with respect to each target variable (predictor and adjacency matrix) separately but not jointly convex. We study the performance of this joint optimization approach through a set of experiments on real and synthetic data. We mention that a byproduct of our investigations is the introduction of a generative model using latent factors for creating artificial data sets. We discuss empirical convergence, implementation issues and performance with respect to each of the two objectives (multivariate prediction and graph inference).

The problem of simultaneous prediction of multivariate time series with graph dependence structure and inference of the graph structure is presented in Section 2. We also present the joint regularization approach for solving this problem. In Section 3 we discuss the algorithm used and the computation of relevant features within this regularization approach. Numerical experiments are discussed and results on synthetic and real data sets are presented in Section 4.

2 Learning Dynamic Graph Features and Edges

2.1 Setup

We first introduce notations for the main objects of interest in the paper:

- a sequence of T undirected weighted graphs with n vertices and their corresponding $n \times n$ adjacency matrices $A_t \in \mathbf{S}_{\geq 0}^n, t \in \{1, 2, ..., T\}$, where $\mathbf{S}_{\geq 0}^n$ denotes the set of $n \times n$ symmetric matrices with nonnegative entries. We assume that, the weights of the edges are nondecreasing functions of time, which in the unweighted case means edges do not disappear but new ones may appear over time.
- a multivariate time series of node features $(X_t)_{1 \le t \le T}$ so that $\forall t, X_t \in \mathbb{R}^{n \times q}$. We denote by $X_t^{(i)} \in \mathbb{R}^q$ the *i*th row of X_t representing the feature vector of node *i* at time *t*. We assume that this series depends on the series of graphs via some function $\omega : \mathbb{R}^{n \times n} \to \mathbb{R}^{n \times q}$, that is, $X_t = \omega(A_t)$.

Given the values of (A_t, X_t) , for $t \in \{1, 2, ..., T\}$, the main goal is to simultaneously predict the future value X_{T+1} and the future adjacency matrix A_{T+1} . To this end, we introduce:

- A descriptor matrix $\Phi_t \in \mathbb{R}^{n \times d}$, $t \in \{1, 2, ..., T\}$, which encodes the past information contained both in the time series $(X_s)_{1 \leq s \leq t}$ and in the sequence of adjacency matrices $(A_s)_{1 \leq s \leq t}$ into a row vector of dimension d for each node.
- A matrix-valued prediction function $f: \mathbb{R}^{n \times d} \to \mathbb{R}^{n \times q}$ which relates the variable X_{t+1} to the past information Φ_t so that X_{t+1} is close to $f(\Phi_t)$. The function f is unknown and has to be estimated from past data. We denote the i-th row component of f with the notation $f^{(i)}$, $i \in \{1, 2, ..., n\}$. The n prediction functions $f^{(i)}$ are assumed to belong to the same Hilbert space \mathcal{H} , equipped with the norm $\|\cdot\|_{\mathcal{H}}$. The norm of $f \in \mathcal{H}^n$ is defined by $\|f\|_{\mathcal{H}^n} := \sqrt{\sum_{i=1}^n \|f^{(i)}\|_{\mathcal{H}}^2}$.

• An unknown matrix $S \in \mathbf{S}_{\geq 0}^n$ whose elements indicate how likely it is that there is a nonzero value at the corresponding position of matrix A_{T+1} . The most likely edges at time T+1 are the ones corresponding to the largest values in S.

2.2 Motivation

The challenge in the prediction problem on real graphs evolving over time is to relate global and slowly varying node features with local and sudden changes of edges. Thus it is reasonable to assume that the evolution of the graph is governed by unobserved latent factors [16, 15], denoted by Ψ_t , which evolve smoothly. For example, in social networks or marketing applications, latent factors could be psychological factors which account for the choices of consumers. The effect of such factors on the graph can then be measured using the adjacency matrix. For instance, the clustering coefficient captures homophilic behavior, the node degree is a good indicator of popularity and the pagerank refers to centrality and influence, the trends can be read in the evolution of node degrees and the tastes are believed to be reflected on the first singular vectors of the adjacency matrix. A possible formulation of the latent factor model sets Ψ_t to be a pair of matrices (U_t, V_t) governing the structure of A_t through $A_t = U_t V_t^{\mathsf{T}}$, where the slow evolution of latent factors is guided by some unobserved underlying mechanism. A rigorous framework would involve a statistical model for the noise and a specific form of dependence of the matrices on their recent history. A formal example will be provided in Section 4. Both the latent factors and their evolution being unobservable, we assume that some relevant and observable features X_t partly capture the information contained in the latent factors.

2.3 Assumptions

The approach studied in this paper relies on the assumption that exploiting simultaneously the *structure of the graph* and the *dynamics of the time* series should improve both prediction and graph learning. Given the information collected over the period t = 1, ..., T, which is contained in the feature matrices Φ_t , we want to learn f and S which satisfy the following assumptions:

• [A1] Low rank. S has low rank. This is a standard assumption in matrix completion problems [18, 5]. The rationale is that the factors U_t, V_t (see Sec. 2.2) should be of small dimension.

- [A2] Graph growth regularity. The growth of the graph exhibits regularity over time, that is, S should be close to the last adjacency matrix A_T . We will use the Frobenius distance between these two matrices as an estimate of the error of the learning process with respect to graph inference.
- [A3] Feature growth regularity. The features of the predicted graph are assumed to be close to the predicted features. As a consequence, we assume that the norm of $f(\Phi_T) \omega(S)$ should be kept small.
- [A4] Stationarity. The dependence mechanism between X_{t+1} and Φ_t can be inferred. Therefore, it is reasonable to assume stationarity of the joint distribution of (X_{t+1}, Φ_t) .
- [A5] Regularity of the predictors over the graph. Neighboring vertices i and j in the graph should correspond to similar prediction functions $f^{(i)}$ and $f^{(j)}$.

2.4 Formulation of the Optimization Problem

We now introduce a regularization based formulation which reflects each of the previous assumptions. To this end, we introduce additional notation. For any matrix M, Tr(M) denotes the trace of M and $\|M\|_F := \sqrt{\text{Tr}(M^T M)}$ denotes the Frobenius norm of M. We also define $\|M\|_* := \sum_{k=1}^n \sigma_k(M)$, the nuclear norm of a matrix M, where $\sigma_k(M)$ denotes the k-th largest singular value of M. We recall that a singular value of matrix M corresponds to the square root of an eigenvalue of $M^T M$ and that the nuclear norm is a convex surrogate of the rank. We also define the matrix $\Delta(f) := (\|f^{(i)} - f^{(j)}\|_{\mathcal{H}}^2)_{i,j=1}^n$. We introduce a convex loss function $\ell : \mathbb{R}^{n \times q} \times \mathbb{R}^{n \times q} \to \mathbb{R}_+$ which measures prediction errors. Given the past history $\{(X_t, A_t) : 1 \le t \le T\}$, the proposed optimization problem is then to minimize over $(f, S) \in \mathcal{H} \times \mathbf{S}_{>0}^n$ the following functional:

$$\mathcal{L}(f,S) := \sum_{t=1}^{T-1} \ell(f(\Phi_t), X_{t+1}) + \frac{\kappa}{2} ||f||_{\mathcal{H}^n}^2 + \ell(f(\Phi_T), \omega(S))$$
$$+\tau ||S||_* + \frac{\nu}{2} ||S - A_T||_F^2 + \lambda \operatorname{Tr}(S^\top \Delta(f)), \qquad (1)$$

where λ , τ and ν are positive regularization parameters. Typical choices of the loss function ℓ are the square loss or the sum of hinge loss errors per

component. Note that using different loss functions allows one to adapt the method to different problems, such as regression or classification.

We now describe each of the separate terms in the above formulation.

- The first term, $J_1(f) = \sum_{t=1}^{T-1} \ell(f(\Phi_t), X_{t+1}) + \frac{\kappa}{2} ||f||_{\mathcal{H}^n}$, corresponds to the prediction task as in standard regularized risk minimization in RKHS for supervised learning.
- The second term, $J_2(S) = \mu ||S||_* + \frac{1}{2} ||S A_T||_F^2$, with $\mu = \tau/\nu$, represents the low-rank matrix denoising objective (static link prediction) and reflects assumptions A1 and A2.
- The third term, $J_3(f,S) = \ell(f(\Phi_T),\omega(S))$, penalizes the difference between the predicted features and the features of the predicted matrix. This term reflects assumption A3.
- The last term, $J_4(f, S) = \text{Tr}(S^{\top}\Delta(f))$ relates the contributions of f and S based on assumption A.

We point out the connection of work on matrix completion [4, 12, 9] to the minimization of $\ell(f(\Phi_T), \omega(S)) + \mu ||S||_*$. Also, in [14], the authors use the minimizer \hat{f} of J_1 and solve the subproblem $\arg \min_S J_2(S) + J_3(S, \hat{f})$.

2.5 Linear Models

For simplicity of presentation, we now consider the case of \mathcal{H} being an RKHS corresponding to a linear kernel. In addition, each $f^{(i)}$ is a linear function represented by a $d \times q$ matrix $W^{(i)}$, so that $f^{(i)}(\Phi_t) = \Phi_t^{(i)}W^{(i)}$. Then we may assume $||f^{(i)}||_{\mathcal{H}} = ||W^{(i)}||_F$ and the graph regularity term can be expressed in terms of the graph Laplacian of S. Recall that the Laplacian is defined as the operator Λ such that $\Lambda(S) = D - S$ where $D = \operatorname{diag}(d_1, \ldots, d_n)$ and $d_i = \sum_{j=1}^n S_{ij}$ are the degrees of the graph. A standard computation gives the following expression for J_4 :

$$\operatorname{Tr}(S^{\top}\Delta(f)) = \sum_{1 \le i,j \le n} S_{ij} \|W^{(i)} - W^{(j)}\|_F^2$$

where $W=(W^{(1)},\dots,W^{(n)})\in\mathbb{R}^{n\times d\times q}$. We use the standard extension of the Frobenius norm to 3-tensors, $\|W\|_F^2=\sqrt{\sum_{i=1}^n\|W^{(i)}\|_F^2}$ and the nota-

tions:

$$Q(W, \Lambda, V) := \sum_{1 \le i, j \le n} \Lambda_{ij} \operatorname{Tr}(W^{(i)^{\top}} V^{(j)}) ,$$
$$\Delta(W) := \left(\|W^{(i)} - W^{(j)}\|_F^2 \right)_{i,j=1}^n .$$

We also assume in this work that the node features are linear functions of adjacency matrices: $\omega(A_t) = A_t \Omega$ for some $\Omega \in \mathbb{R}^{n \times q}$. Note that degree, inter/intra cluster degrees and projection onto specific linear subspaces are such features, but not clustering coefficients or statistics on path lengths. We later detail how we chose Ω in our experiments.

Working with a known and static graph. This problem was considered in [6] showing interestingly that feature prediction can be improved when using graph structure in the regularized optimization formulation. In their work, the graph penalty (our J_4) takes the form:

$$\sum_{1 \le i,j \le n} S_{ij} \|W^{(i)} - W^{(j)}\|_{1} ,$$

where S is fixed and we can see that an L_1 norm is used for this penalty term instead of L_2 in our case. Moreover, their approach allows negative interactions between vertices. This indicates that there are clearly many different variants to explore.

3 Learning Algorithm

In this section we address several issues related to the optimization problem (2.4) and discuss the learning algorithm we will use. The algorithm is a standard *projected gradient* method, which projects on the constraint set \mathcal{E} defined below. The main challenges are the nonconvexity of J_4 and the nondifferentiability of J_2 in the objective.

3.1 Convexity

In the experiments conducted, we have learned linear functions $f^{(i)}$. In this case, the term $\sum_{1 \leq i,j \leq n} S_{ij} ||W^{(i)} - W^{(i)}||_F^2$ is convex with respect to one of the variables, if the other is fixed, but not jointly convex with respect to (W,S). In order to ensure convergence of an optimization algorithm to a desired minimum, we determine a domain where \mathcal{L} is convex. We use intuition

from what happens with a similar optimization problem in dimension 2. Indeed, consider the minimization of $(w,s) \to sw^2 + \alpha s^2 + \beta w^2$ which is the simplified functional in the degenerate case where n=1. The eigenvectors of the Hessian are positive iff $w^2 \leq \alpha \beta$, so this condition defines the convexity region. By adding the quadratic terms in S and W to $Q(W, \Lambda(S), W)$, let us define

$$\Psi(S, W) := \frac{\kappa}{2} \|W\|_F^2 + \frac{\nu}{2} \|S - A_T\|_F^2 + \lambda Q(W, \Lambda(S), W).$$

If we suppose that the entries of $Z=S-A_T$ are nonnegative, then the minimizer of Ψ is the trivial solution $S=A_T, W=0$. In fact,

$$\Psi(S, W) = \frac{\kappa}{2} \|W\|_F^2 + \frac{\nu}{2} \|Z\|_F^2 + \lambda \left(Q(W, \Lambda(A_T), W) + Q(W, \Lambda(Z), W) \right) \ge 0 \quad (2)$$

and for $Z = 0, W = 0, \Psi(S, W) = 0.$

We prove that \mathcal{L} is convex over a set \mathcal{E} around the minimizer of Ψ and we will ensure henceforth that the descent algorithm takes place inside this convex domain.

Proposition 1 The function Ψ is convex in the interior of the set:

$$\mathcal{E} = \left\{ S \in \mathbf{S}_{\geq 0}^n, W \in \mathbb{R}^{n \times d \times q} \mid \|W\|_F \le \frac{\sqrt{\nu \kappa}}{2\lambda(\sqrt{n} + 1)} \right\}.$$

Proof. We introduce the slack variable $Z = S - A_T$ and isolate the quadratic part of $\Psi(Z_0 + Z, W_0 + W)$ for some (Z_0, W_0) :

$$R(Z,W) := \frac{\nu}{2} \|Z\|_F^2 + \frac{\kappa}{2} \|W\|_F^2 + \lambda \Big(2Q(W, \Lambda(Z), W_0) + Q(W, \Lambda(A_T + Z_0), W) \Big).$$

Thanks to Cauchy-Schwarz and the basic norm property $||AB||_F \leq ||A||_F ||B||_F$,

$$Q(W, \Lambda(Z), W_0) \ge -\|W\|_F \|\Lambda(Z)\|_F \|W_0\|_F.$$

We have $\Lambda(Z) = D - Z$. We get, again by Cauchy-Schwarz

$$||D||_F^2 = \sum_{i=1}^n (\sum_{j=1}^n Z_{i,j})^2 \le \sum_{i=1}^n n \sum_{j=1}^n Z_{i,j}^2 = n ||Z||_F^2$$

and therefore

$$||\Lambda(Z)||_F = ||D - Z||_F \le ||D||_F + ||Z||_F \le (\sqrt{n} + 1)||Z||_F.$$

On the other hand $Q(W, \Lambda(A_T + Z_0), W) \ge 0$, so

$$R(Z,W) \ge \frac{\nu}{2} \|Z\|_F^2 + \frac{\kappa}{2} \|W\|_F^2 - 2\lambda(\sqrt{n} + 1) \|W\|_F \|Z\|_F \|W_0\|_F.$$

Letting $z = ||Z||_F$, $w = ||W||_F$, $w_0 = ||W_0||_F$, we have

$$R(Z, W) \ge \frac{\nu}{2} z^2 + \frac{\kappa}{2} w^2 - 2\lambda(\sqrt{n} + 1)wzw_0$$
.

 $\psi_{w_0}: (z,w) \mapsto \frac{\nu}{2}z^2 + \frac{\kappa}{2}w^2 - 2\lambda(\sqrt{n}+1)wzw_0$ is a quadratic form. Therefore R(Z,W) is always nonnegative if ψ_{w_0} is positive semidefinite positive. That is, if $\nu + \kappa \geq 0$ (always true) and $\nu\kappa - 4\lambda^2(\sqrt{n}+1)^2w_0^2 \geq 0$, and this completes the proof.

3.2 Smoothing the Nuclear Norm

For optimization purposes, we chose to replace the nonsmooth nuclear norm term by a smooth approximation, in the spirit of the optimization literature, for example [13],

$$g_{\eta}(S) = \max_{Z} \left\{ \langle S, Z \rangle - \frac{\eta}{2} ||Z||_{F}^{2} ||\sigma_{1}(Z) \leq 1 \right\}.$$

Each of these parameterized functions is a lower bound that approaches $||S||_*$ as $\eta \to 0^+$, while being differentiable. Differentiability is due to $-g_{\eta}$ being related to a Moreau envelope [11] and specifically to the squared distance of $\frac{1}{\eta}S$ from the unit ball of the spectral norm.

3.3 Gradient Evaluation

In the case of a differentiable loss ℓ , we explicitly compute $\nabla \mathcal{L}(W,S)$ as

$$\frac{\partial \mathcal{L}}{\partial W}(W, S) = \sum_{t=1}^{T-1} \nabla_W \ell(W\Phi_t, X_t) +$$

$$\nabla_W \ell(W\Phi_T, S\Omega) + \kappa W + 2\lambda \left(\sum_{j=1}^n \Lambda_{ij}(S)W^{(j)}\right)_{i=1}^n$$

and

$$\frac{\partial \mathcal{L}}{\partial S}(W, S) = \nabla_{S} \ell(W \Phi_{T}, S\Omega) + \lambda \Delta(W) + \nu(S - A_{T}) + \tau U(S) \operatorname{Diag}\left(\min\left(1, \frac{\sigma_{i}(S)}{\eta}\right)\right) V(S)^{\mathsf{T}},$$

where $S = U(S) \operatorname{Diag}(\sigma_1(S), \dots, \sigma_n(S)) V(S)^{\top}$ is a singular value decomposition of S. In the case of squared loss,

$$\ell(W\Phi_t, X_t) = \sum_{i=1}^n \|X_t^{(i)} - \Phi_t^{(i)} W^{(i)}\|_2^2.$$

Therefore:

$$\nabla_{S}\ell(W\Phi_{T}, S\Omega) = \left(\left(S\Omega\right)^{(i)} - \Phi_{T}^{(i)}W^{(i)}\Omega^{\top}\right)_{i=1}^{n},$$

$$\nabla_{W}\ell(W\Phi_{t}, X_{t}) = \left(\Phi_{t}^{(i)^{\top}}(\Phi_{t}^{(i)}W^{(i)} - X_{t}^{(i)})\right)_{i=1}^{n}.$$

4 Numerical Experiments

4.1 Description of Data Sets

4.1.1 Synthetic Data

We introduce a generative model for the graph sequence and the node-related features:

$$\begin{cases} X_t = A_t \Omega \\ A_t = U_t V_t^\top + z_t \\ \forall i \begin{cases} U_t^{(i)} = U_{t-1}^{(i)} + h(U_{t-1}^{(i)}) + u_{t,i} \\ V_t^{(i)} = V_{t-1}^{(i)} + h(V_{t-1}^{(i)}) + v_{t,i} \end{cases} \end{cases}$$

where $u_{t,i}, v_{t,i}$ are multivariate Gaussian vectors $\mathcal{N}(0, \delta^2 I_r)$ in \mathbb{R}^r , and the entries of z_t are independently drawn from a centered Gaussian with variance σ^2 . We also set

$$h(x) = \epsilon \left(e^{\frac{-\|x - v_1\|^2}{\sigma_1^2}} (x - v_1) + e^{\frac{-\|x - v_2\|}{\sigma_2}} (x - v_2) \right),$$

where vectors $v_1, v_2 \in \mathbb{R}^r_+$ are chosen randomly and $\epsilon, \sigma_1, \sigma_2$ are positive constants. We use here n = 100, r = 4, T = 60 and the entries of $U_0, V_0 \in \mathbb{R}^{n \times r}$

are drawn according to a uniform distribution in (0,1). Note that such data by construction fulfill the required assumptions, and that the nonlinearity of the smooth function h is what makes the contribution of the Laplacian term $J_4(f,S)$ non-trivial, we highlight this phenomenon in Figure 1.

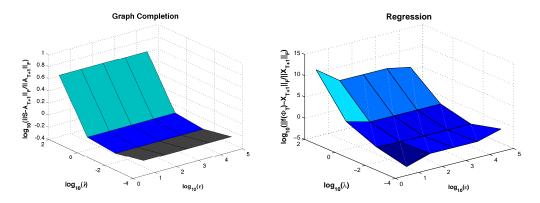


Figure 1: Prediction accuracy as a function of λ (predictors' regularity on the graph) and ϵ (nonlinearity of latent factors growth).

4.1.2 Marketing Data

We use data from several months of book purchase history of a major ecommerce website to evaluate our method. We tested our method on the task of predicting the cross-selling and the sales volumes on two types of data sets: (1) the top-sold items and (2) at an aggregated level among different product categories.

1. Best-sellers For our test we selected a set of 300 top-sold books, and we aimed at predicting the future sales of those books, and also predicting the cross-sales, or co-purchases. We consider the set of products being the vertices of the dynamic graph sequence, and construct the 'co-purchase' graph as follows. At time t one can represent the state of market by a binary #users × #items matrix M_t where entry (i, j) is non-zero if user i has purchased item j so far. We define the co-purchase matrix A_t = M_t^T M_t as the adjacency matrix of the weighted graph. The co-purchase graph consists of the products as nodes and the weights of edges are given by the number of co-purchases of two products. We use 32 weeks (8 months) of observation for learning, and predict the evolution of sales volumes of the books over the next 5

weeks. The target feature, namely the sales volume is obtained using the bipartite graph of users and products. The degrees of items in the bipartite graph represent the sales volumes. Our algorithm outputs simultaneously predictions of the top-selling items and a matrix S predicting the cross-sellings.

2. Categories Predicting the sales and co-sales at the aggregated level of product categories is relevant for marketers and supply-chain managers. We collected data from n = 195 book categories over a history of T = 74 week periods. The entry (i, j) of A_t counts the number of users having purchased items of category i at time t-1 who purchased at least an item from category j at time t. The features of interest are the weekly sales of each category and the number of new users in each category.

4.2 Features and Descriptors

We recall that our setup consists of a regression-type model for the prediction of (X_{T+1}, A_{T+1}) given the available information at time T which includes the previous adjacency matrices A_1, \ldots, A_T of size $n \times n$ and the node-related feature matrices X_1, \ldots, X_T of size $n \times q$. Our approach involves a Φ_T of size $n \times d$ called descriptor matrix which encompasses the information contained in the past realizations (X_t, A_t) for $t \leq T$.

4.2.1 Node Features

The feature matrix X_t contains two types of node features:

- 1. Features whose predicted values are of direct interest. For instance, the volume of sales or popularity of an item, represented by the degree of the related vertex in the purchase graph, or the growth rates of the degree as an indicator of the penetration level [17] of the item in the market.
- 2. Features which are believed to measure relevant underlying factors that govern graph structure. For instance the clustering coefficient [20] taken as a proxy for homophily is not of direct interest, but is a useful quantity for inferring the future connections of a node.

For simplicity, we set here $X_t = A_t \Omega$ where Ω is an $n \times q$ matrix. We assume that the columns of Ω contain (i) the constant vector $\mathbf{1}_n$ (since $A_t \mathbf{1}_n$ is the vector of node degrees), (ii) the indicator vectors of several clusters which

we identified in the graph and (iii) the top k = 5 eigenvectors of A_t which are believed to capture most of the structure of A_t .

4.2.2 Node Descriptors

The descriptors contain the features and also other quantities that are useful for predicting the features. In the present work, we made the following choice:

$$\Phi_t = (A_t \Omega, (A_t - A_{t-1})\Omega, (A_t - 2A_{t-1} + A_{t-2})\Omega)$$
$$= (X_t, X_t - X_{t-1}, X_t - 2X_{t-1} + X_{t-2})$$

which is a matrix of size $n \times 3q$ representing features, their velocity and acceleration rates. We remark that time-related statistics of the recent history of $(X_t)_{t \leq T}$ (such as moving averages, residual variance), or other quantities accounting for alternative representations of the time series (such as Fourier or wavelet coefficients, polynomial approximation) could also be included.

4.3 Evaluation Metrics

Various prediction tasks could be studied in our setup with specific criteria for optimization and evaluation. We focus here on regression (for feature prediction) and graph completion (for link prediction), but classification of vertices may also be a useful task.

Regression. When the effective value of an asset in the future is of interest, a squared error metric is appropriate, leading to the use of $\|X_{T+1} - f(\Phi_T)\|_F$ for evalution and $l(f(\Phi_t), X_t) = \|X_{t+1} - f(\Phi_t)\|_F^2$ in the objective. The reported values are the relative errors $\frac{\|X_{T+1} - f(\Phi_T)\|_F}{\|X_{T+1}\|_F}$.

Graph Completion. Our method aims simultaneously at the prediction

Graph Completion. Our method aims simultaneousely at the prediction of X_{T+1} and A_{T+1} . We measure the quality of prediction of A_{T+1} by $||S - A_{T+1}||_F$, and the relative values reported are $\frac{||A_{T+1} - S||_F}{||A_{T+1}||_F}$.

Classification. Specific patterns may appear in the time series X_t which we may wish to predict. For instance, predicting top-selling items in a given market is definitely of interest. This problem can be formulated as follows: the matrix X_t represents the sales volumes over a market (each component corresponds to a product), and we assign a ± 1 label depending on the order of magnitude of the increase of sales volumes over a specific time window.

4.4 Comparison with Related Methods

Since the problem which we address (simultaneous graph completion and feature prediction) is a novel one, there are no direct competing methods to compare with. Thus, in our experiments we have considered the following baselines which address the most similar learning problems:

- 1. Ridge Regression. Learning the function $\hat{f} := \arg\min_{\mathcal{H}} J_1$ through the minimization of the penalized regression $J_1(f) = \sum_{t=1}^{T-1} \ell(f(\Phi_t), X_{t+1}) + \frac{\kappa}{2} ||f||_{\mathcal{H}^n}^2$ leads us to a prediction function \hat{f} which we use for predicting X_{T+1} by $f(\Phi_T)$. This approach addresses the prediction of node features as independent time series. It does not take the graph sequence into account and does not predict the future graph.
- 2. Rank-free prediction. We used our formulation with the choice of $\tau=0$. In this case, the loss function ignores the low-rank hypothesis in the graph structure. The prediction process only relies on the dynamics of graph features.
- 3. Graph completion through shrinkage. We also considered a method only dedicated to link prediction and matrix denoising, ignoring the feature prediction part of the problem. This method only uses the observation A_T as a noisy version of A_{T+1} , and aims at predicting A_{T+1} given only this observation and the low rank hypothesis. We solve the problem $\min_S J_2(S)$ where $J_2(S) = \frac{1}{2} ||S A_T||_F^2 + \mu ||S||_*$ to estimate the adjacency matrix as a low rank approximation of A_T . The solution of this problem is obtained by singular value shrinkage [4], namely if $A_T = U \operatorname{diag}(\sigma_1(A_T), \ldots, \sigma_n(A_T))V^{\top}$ is the singular value decomposition of A_T , then the image of A_T by a shrinkage of parameter μ is given by $D_{\mu}(A_T) = U \operatorname{diag}(\sigma_i(A_T) \mu)_+ V^{\top}$.

4.5 Results

We implemented a first-order gradient descent algorithm with projections on the convex set \mathcal{E} at each iteration. In Figure 2 we show how cross validation can be done first on the two standard objectives of regression and graph completion (matrix denoising) taken separately for setting the values of $\mu = \frac{\tau}{\nu}$ and κ efficiently. This reduces the number of smoothing parameters, and thanks to experiments on synthetic data we studied the dependence of optimal regularization parameters $\kappa, \tau, \nu, \lambda$ on the model parameters r, σ, ϵ . We illustrate in Figure 3 the efficiency of our algorithm for minimizing different

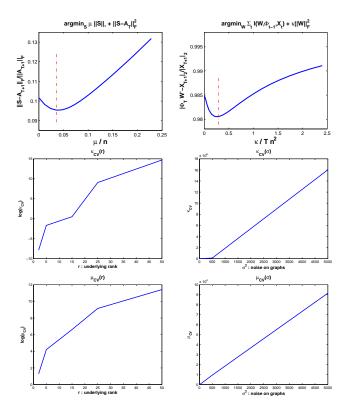


Figure 2: Cross-validation for μ and κ on the two standard objectives of regression and graph completion (matrix denoising) taken separately (top 2 plots) and dependence of constants chosen by cross-validation $\mu_{\rm CV}, \kappa_{\rm CV}$ on generative model constants r, σ .

objective terms despite competing effects at the starting point of optimization. We emphasize the decrease of validation error as a function of the number of iterations. Finally, in Figure 4, we report results of experiments on marketing data sets, which show that, with appropriately chosen hyperparameters, our approach outperforms standard competing baselines. We also observe that on the top-selling items data set, the nuclear norm term does not add to prediction accuracy compared to the rank-free $\tau=0$ case. Further study is needed to better understand whether this effect is due to the possible presense of high noise or to the higher efficiency of predicted features in predicting the future graph.

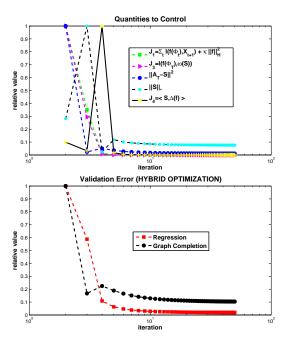


Figure 3: Numerical behavior of different terms of the risk functional and validation error versus iteration count of a gradient descent algorithm (synthetic data).

5 Conclusion

We have studied the benefits of joint regularization for simultaneously predicting time series which are related, as well as the graph representing these relations. The main goal has been to investigate the feasibility as well as the advantages of such a simultaneous estimation procedure. Although the corresponding optimization problem is not convex, we identified a convexity domain which in some cases will contain the optimal solution. The experimental results indicate that exploring such joint regularization and learning problems can greatly improve the predictive performance of these methods. Clearly, an important question is to determine conditions under which the global optimum is attained within the convexity domain of the optimization problem (2.4). A key contribution of this work is to show the potential of this hybrid approach, as well as to propose a practical algorithm for solving the learning problem. Moreover, given the promising empirical results, future improvements in the optimization methodology may lead to even further improvements in terms of predictive performance.

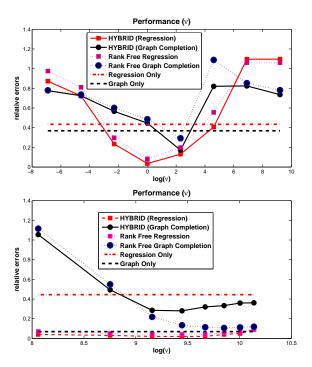


Figure 4: Performance on real data for different values of ν , where we have set $\tau = \nu \mu_{\rm CV}$ for a value of $\mu_{\rm CV}$ fixed in advance by cros-validation. Book categories sales volumes and cross-selling on the top and best-selling items on the bottom.

References

- [1] Jake Abernethy, Francis Bach, Theodoros Evgeniou, and Jean-Philippe Vert. A new approach to collaborative filtering: operator estimation with spectral regularization. *JMLR*, 10:803–826, 2009.
- [2] Andreas Argyriou, Massimiliano Pontil, Charles Micchelli, and Yiming Ying. A spectral regularization framework for multi-task structure learning. *Proceedings of Neural Information Processing Systems* (NIPS), 2007.
- [3] Leo Breiman and Jerome H. Friedman. Predicting multivariate responses in multiple linear regression. *Journal of the Royal Statistical Society (JRSS): Series B (Statistical Methodology)*, 59:3–54, 1997.

Data	Marketing(BS)		Marketing(cat.)		Synthetic	
Method \ Res.	Feature	Graph	Feature	Graph	Feature	Graph
Hybrid	0.06	0.33	0.08	0.31	$0.13 \pm .002$	$0.21 \pm .003$
$\lambda = 0: [14]$	0.12	0.32	0.09	0.34	$0.16 \pm .003$	$0.22 \pm .004$
$\tau = 0$: Rank-free	0.08	0.39	0.06	0.24	$0.19 \pm .008$	$0.24 \pm .01$
Regression Only	-	0.42	-	0.17	-	$0.23 \pm .003$
Graph Only	0.40	-	.31	-	$0.21 \pm .01$	-

Table 1: Comparison of our method with other methods on real marketing data and synthetic data sets. Hyper-parameters set after a cross validation step on past data. Synthetic data results are reported with bootstrap standard type errors computed after 50 replications. Feature prediction relative residuals $\frac{\|X_{T+1}-f(\Phi_T)\|_2}{\|X_{T+1}\|_2}$ and Graph prediction relative residuals $\frac{\|A_{T+1}-S\|_F}{\|A_{T+1}\|_F}$ reported.

- [4] Emmanuel J. Candès and Terence Tao. A singular value thresholding algorithm for matrix completion. SIAM Journal on Optimization, 20(4):1956–1982, 2008.
- [5] Emmanuel J. Candès and Terence Tao. The power of convex relaxation: Near-optimal matrix completion. *IEEE Transactions on Information Theory*, 56(5), 2009.
- [6] Xi Chen, Seyoung Kim, Qihang Lin, Jaime G. Carbonell, and Eric P. Xing. Graph-structured multi-task regression and an efficient optimization method for general fused lasso. *arXiv*, 2010.
- [7] Theodoros Evgeniou, Charles A. Micchelli, and Massimiliano Pontil. Learning multiple tasks with kernel methods. *Journal of Machine Learning Research*, 6:615–637, 2005.
- [8] Mladen Kolar and Eric P. Xing. On time varying undirected graphs. in Proceedings of the 14th International Conference on Artifical Intelligence and Statistics AISTATS, 2011.
- [9] Vladimir Koltchinskii, Karim Lounici, and Alexandre Tsybakov. Nuclear norm penalization and optimal rates for noisy matrix completion. *Annals of Statistics*, 2011.
- [10] David Liben-Nowell and Jon Kleinberg. The link-prediction problem for social networks. *Journal of the American Society for Information Science and Technology*, 58(7):1019–1031, 2007.

- [11] Jean-Jacques Moreau. Proximité et dualité dans un espace hilbertien. Bulletin de la Société Mathématique de France, 93(2):273–299, 1965.
- [12] Sahand Negahban and Martin J. Wainwright. Estimation of (near) low-rank matrices with noise and high-dimensional scaling. *Annals of Statistics*, 39:1069–1097, 2011.
- [13] Yurii Nesterov. Smooth minimization of non-smooth functions. *Mathematical Programming*, 103(1):127–152, 2005.
- [14] Emile Richard, Nicolas Baskiotis, Theodoros Evgeniou, and Nicolas Vayatis. Link discovery using graph feature tracking. *Proceedings of Neural Information Processing Systems (NIPS)*, 2010.
- [15] Purnamrita Sarkar, Deepayan Chakrabarti, and Andrew W. Moore. Theoretical justifiation of popular link prediction heuristics. In Proceedings of COLT, 2010.
- [16] Purnamrita Sarkar and Andrew Moore. Dynamic social network analysis using latent space models. In Y. Weiss, B. Schölkopf, and J. Platt, editors, *Proceedings of Neural Information Processing Systems* 18, pages 1145–1152. MIT Press, Cambridge, MA, 2006.
- [17] Ashish Sood, Gareth M. James, and Gerard J. Tellis. Functional regression: A new model for predicting market penetration of new products. Marketing Science, 28(1):36–51, 2009.
- [18] Nathan Srebro, Jason D. M. Rennie, and Tommi S. Jaakkola. Maximum-margin matrix factorization. In Lawrence K. Saul, Yair Weiss, and Léon Bottou, editors, in Proceedings of Neural Information Processing Systems 17, pages 1329–1336. MIT Press, Cambridge, MA, 2005.
- [19] Ruey S. Tsay. Analysis of Financial Time Series. Wiley-Interscience; 3rd edition, 2005.
- [20] Ducan J. Watts and Steven H. Strogatz. Collective dynamics of small-world networks. *Nature*, 393:440–442, 1998.

